



Git

A great development tool that works
nicely with Perforce



*Jordan Zimmerman
Senior Platform Engineer
Platform Engineering*

Disclaimer!

I am not a Git Expert*

Yes, you will be able to stump the speaker if you try



**But I'm willing to
help!**

**Google is your
friend!**



**Though I have used it extensively for more than a year*

Disclaimer #2

- I'm not trying to sell anyone on Git
- I like it – you might not
- Use what you like



Agenda



Quick History for the Curious



Git Primer



But What About Perforce?



Git and Perforce at the Same Time!



Git Tips and Tricks



Demos!



Quick History for the Curious

BitKeeper originally used as SCM for Linux

Free use of BitKeeper was revoked

“I hate CVS with a passion”

-Linus Torvalds

“I see Subversion as being the most pointless project ever started”

-Linus Torvalds



Given the above, guess what Torvalds did?

P.S. I don't agree with Linus

Torvalds had several criteria for an SCM system:

- Take CVS as an example of what not to do; if in doubt, make the exact opposite decision
- Support a distributed, BitKeeper-like workflow
- Very strong safeguards against corruption, either accidental or malicious
- Very high performance



*No SCM met all of these –
so...*

Enter *Git*

- A distributed version control system focused on speed, effectiveness* and real-world usability on large projects

Strong support for non-linear development

Distributed development

Efficient handling of large projects

Cryptographic authentication of history

Toolkit
design

“I’m an egotistical bastard and I name all my projects after myself. First Linux, now Git.”

-Linus Torvalds



*I have no idea what this word means – I got it from the Git website

Git Primer

- Git is different than your Father's SCM
- Local repository
- No need to checkout files
- No odd client view mappings
- Very branch friendly
- Staging area
- Github



Git Primer

- Git is **very** different than your Father's SCM
- Local repository
- No need to checkout files
- No odd client view mappings
- Very branch friendly
- Staging area
- Github



Git Hashes

- All git repo objects identified by a SHA-1 hash
 - Directories
 - Files
 - Commits
- Any commit/object can be ID'd by its hash or beginning portion of its hash

```
SHA: ccc90efddcdc760a108b7b9d711b6a87d004f066
Author: Jordan Zimmerman <jzimmerman@netflix.com>
Date: Tue Jul 19 2011 09:34:22 GMT-0700 (PDT)
Subject: activated all tests
Refs: async
Parent: 0ae6893a7e0139e3348cb416b4ec43090128e82c
```



Git Primer – cont'd

- The next few slides are copied from:

Git the Basics
by Bart Trojanowski

<http://edgyu.excess.org/git-tutorial/2008-07-09/intro-to-git.pdf>



Git Primer - Concepts

- Source Control Management
 - Track changes to files
 - repository / database of changes
 - working directory / current state
- Centralized SCM
 - server: single database
 - client: working directory & state

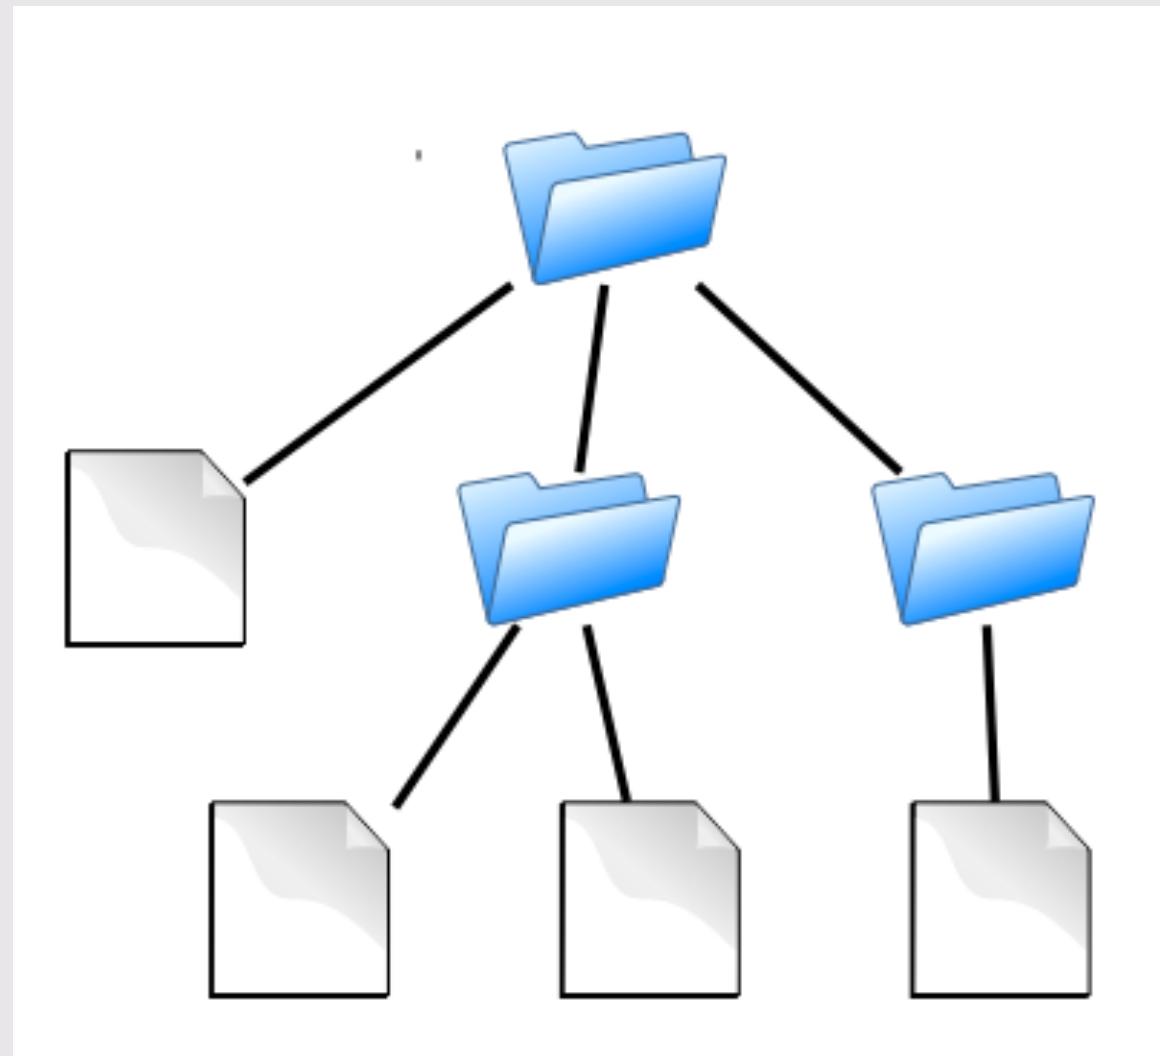


- Decentralized SCM

- Anyone can be a server

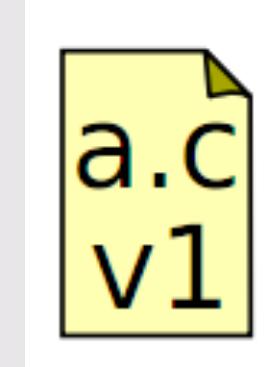
Git Primer – SCM Components

- Working tree
 - Directories
 - Files



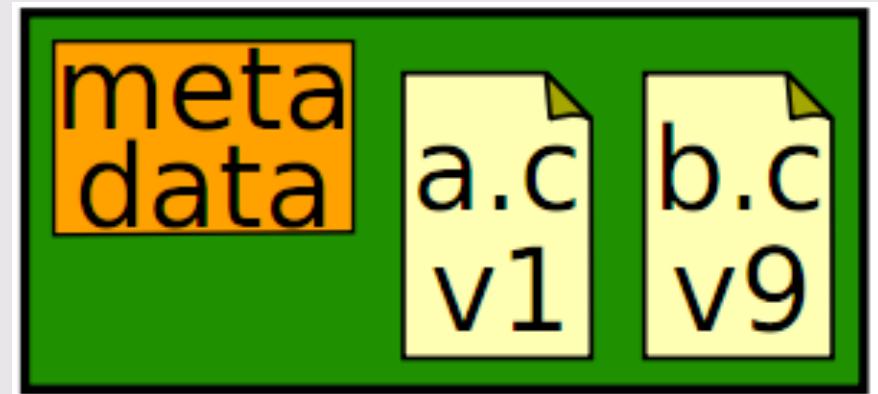
Git Primer – SCM Components

- Repository contents
 - Files



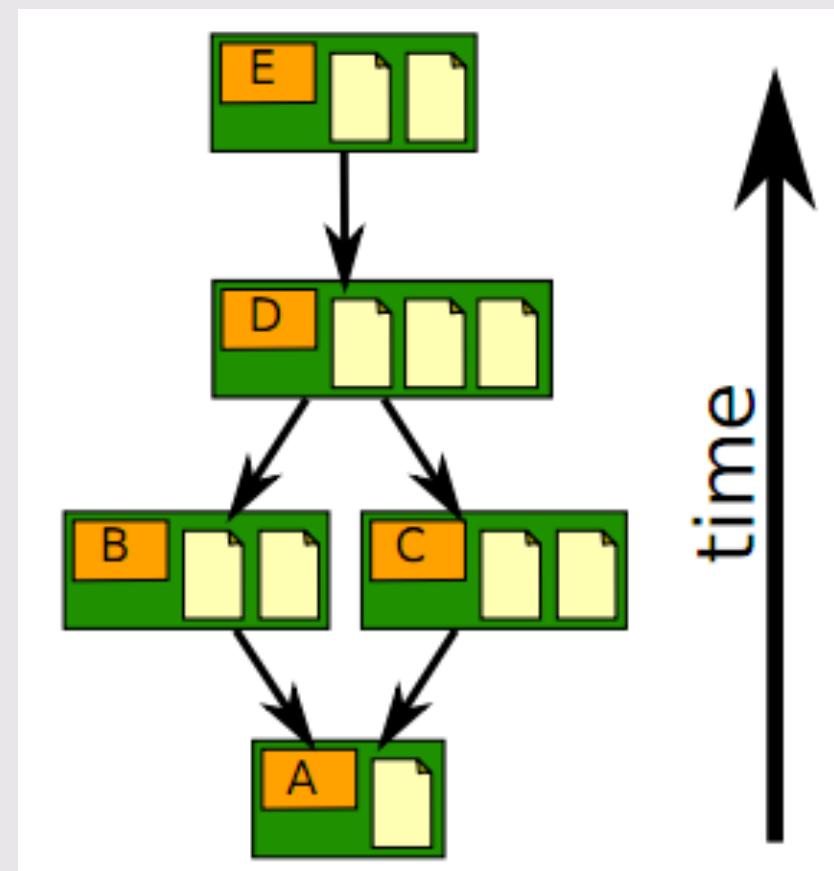
Git Primer – SCM Components

- Repository contents
 - Files
 - Commits



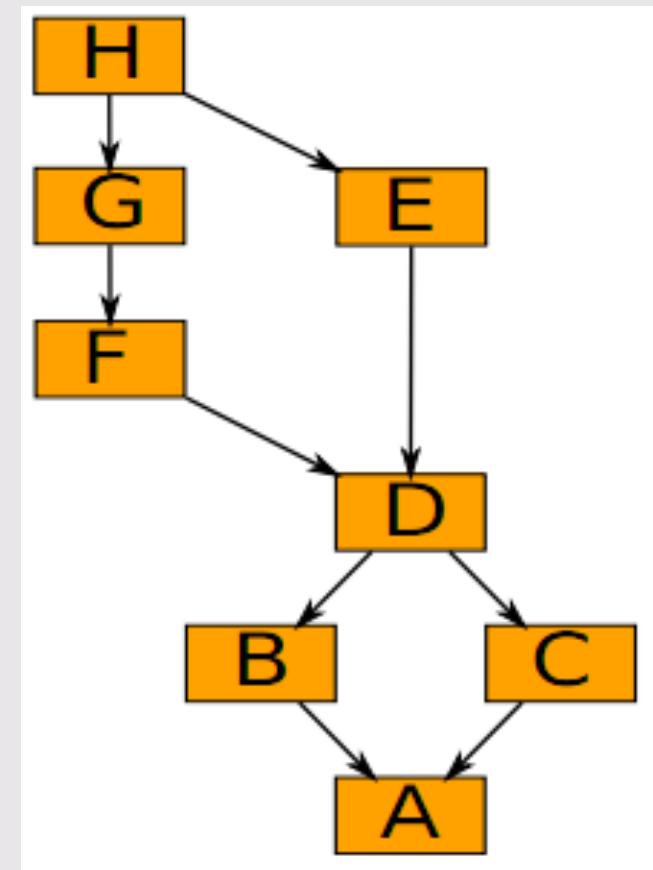
Git Primer – SCM Components

- Repository contents
 - Files
 - Commits
 - Ancestry



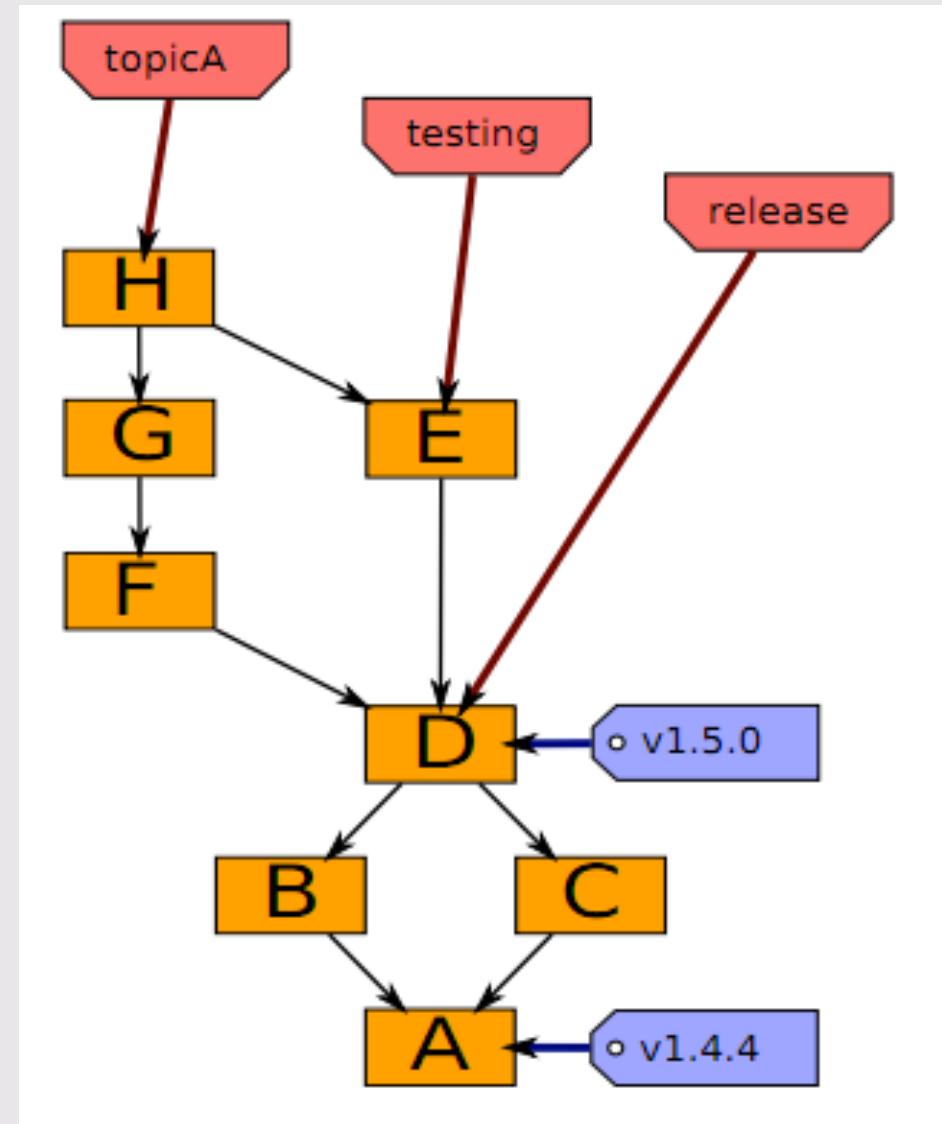
Git Primer – SCM Components

- Directed Acyclic Graph (DAG)



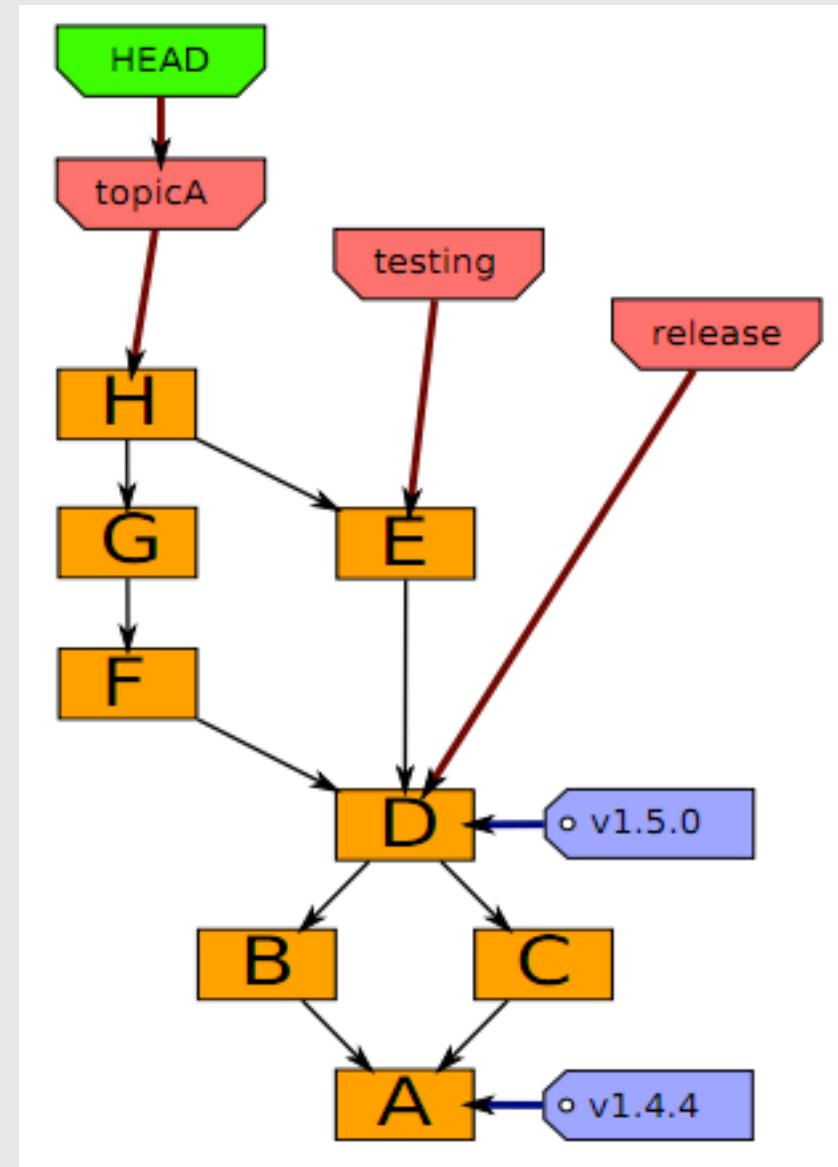
Git Primer – SCM Components

- References
 - Tags
 - Branches



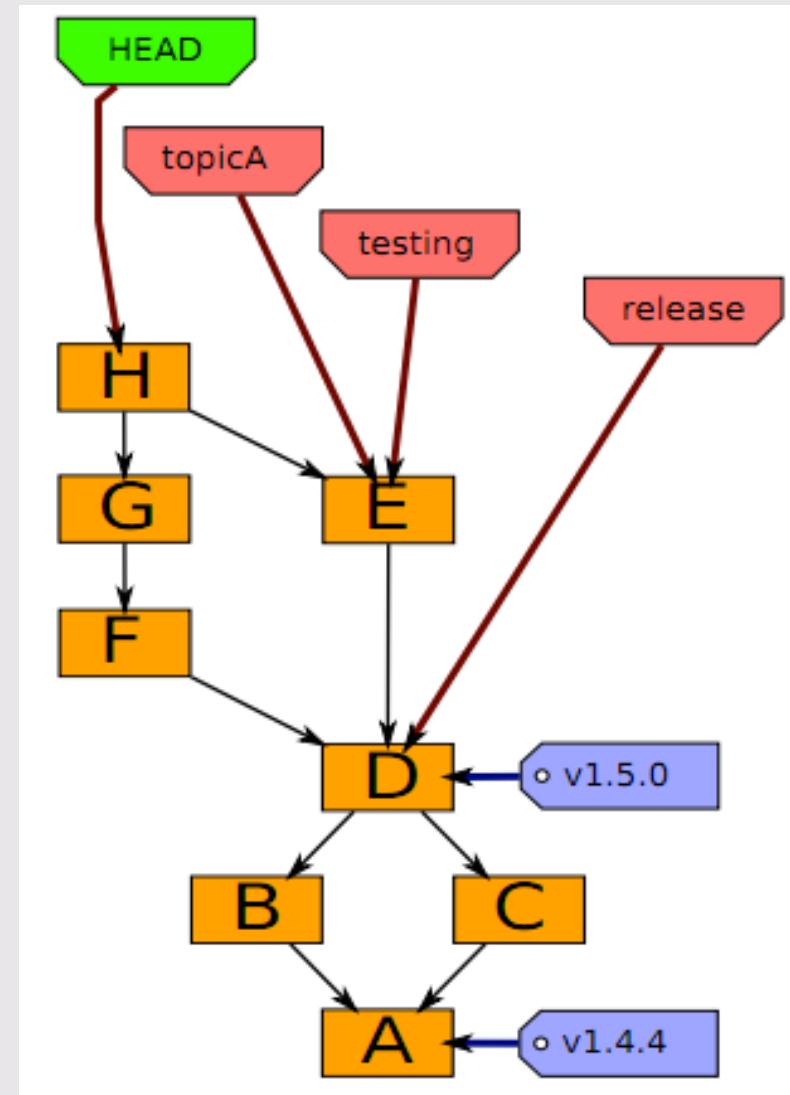
Git Primer – SCM Components

- HEAD
 - Current checkout
 - Points to branch



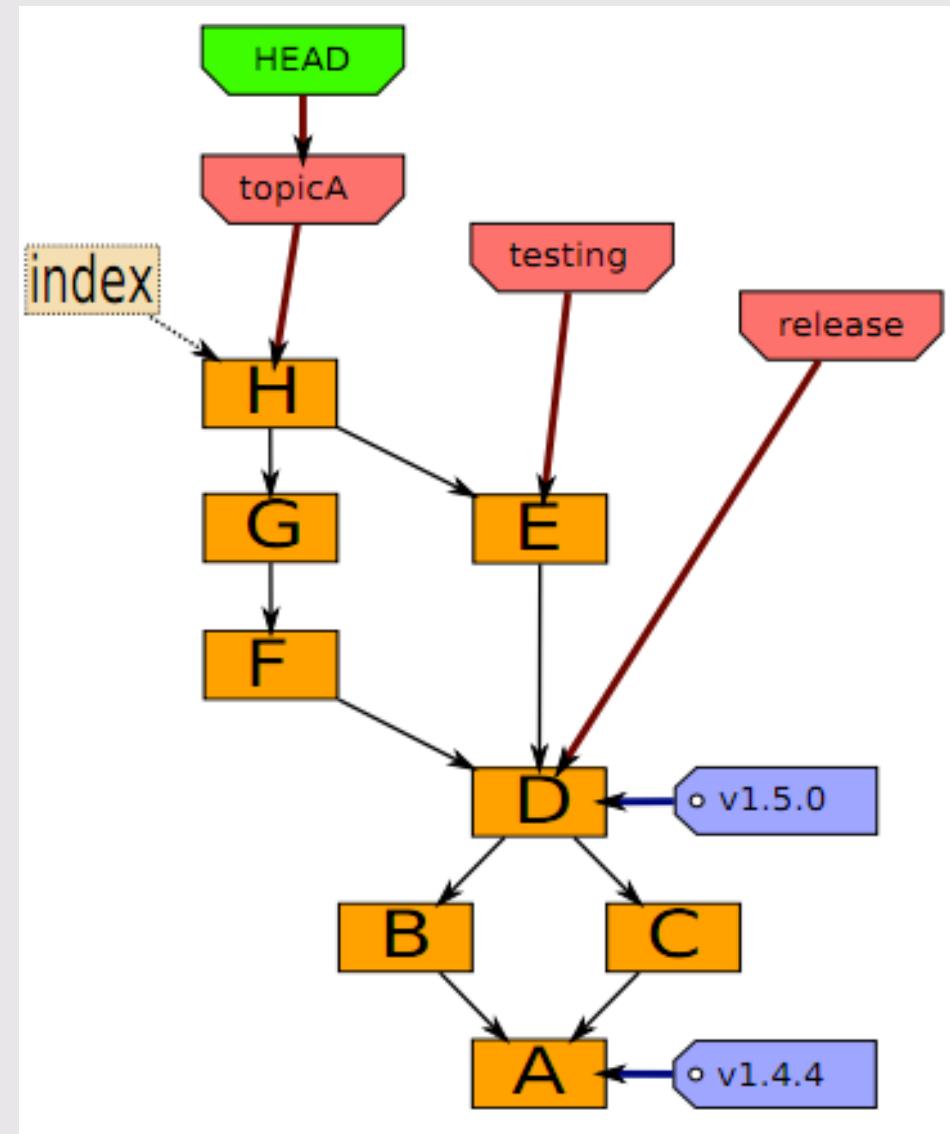
Git Primer – SCM Components

- HEAD
 - Current checkout
 - Points to branch
 - Sometimes detached



Git Primer – SCM Components

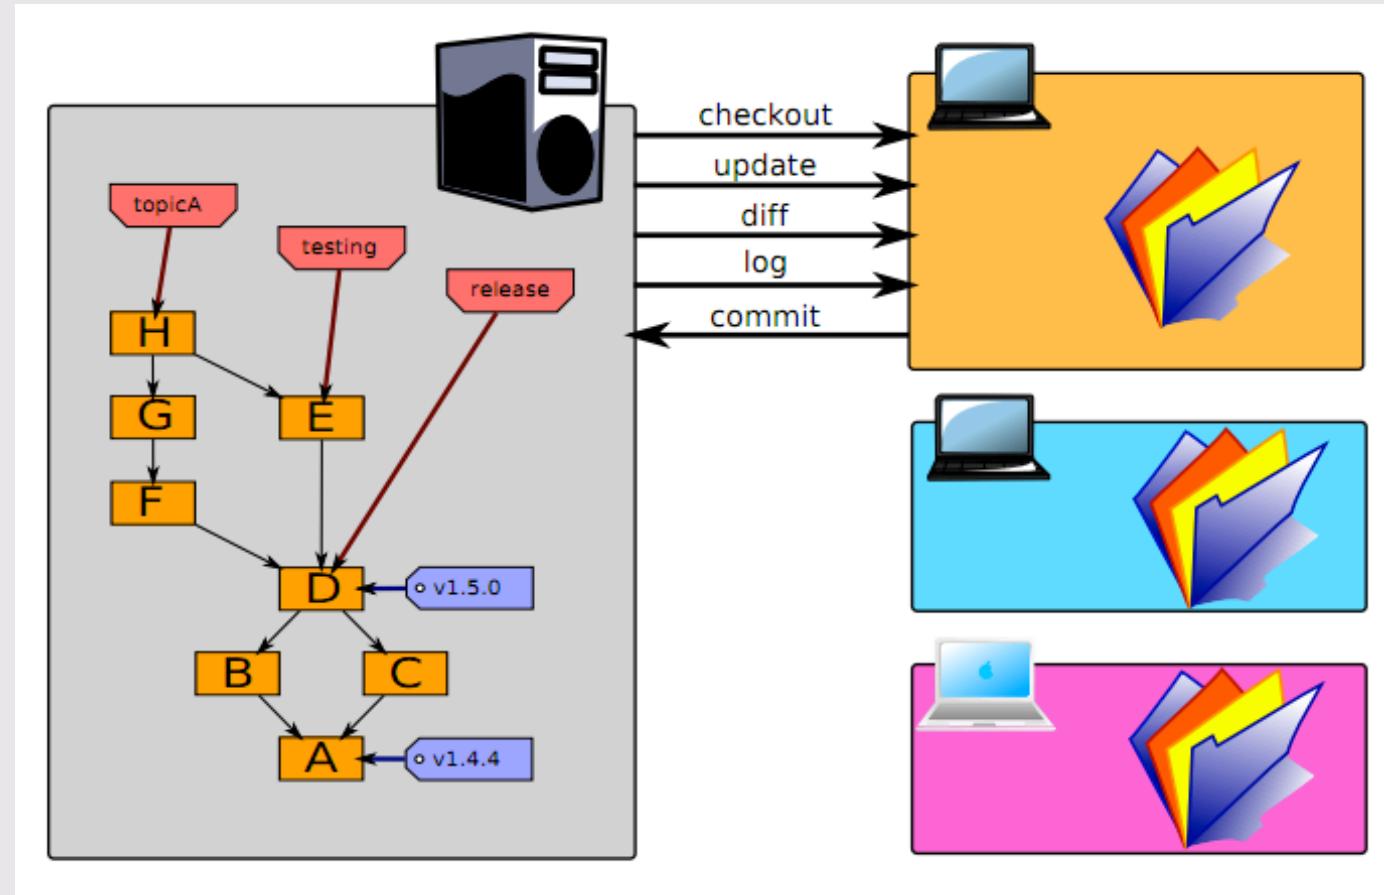
- Index
 - Staging area
 - What is to be committed



Git Primer – Centralized SCM

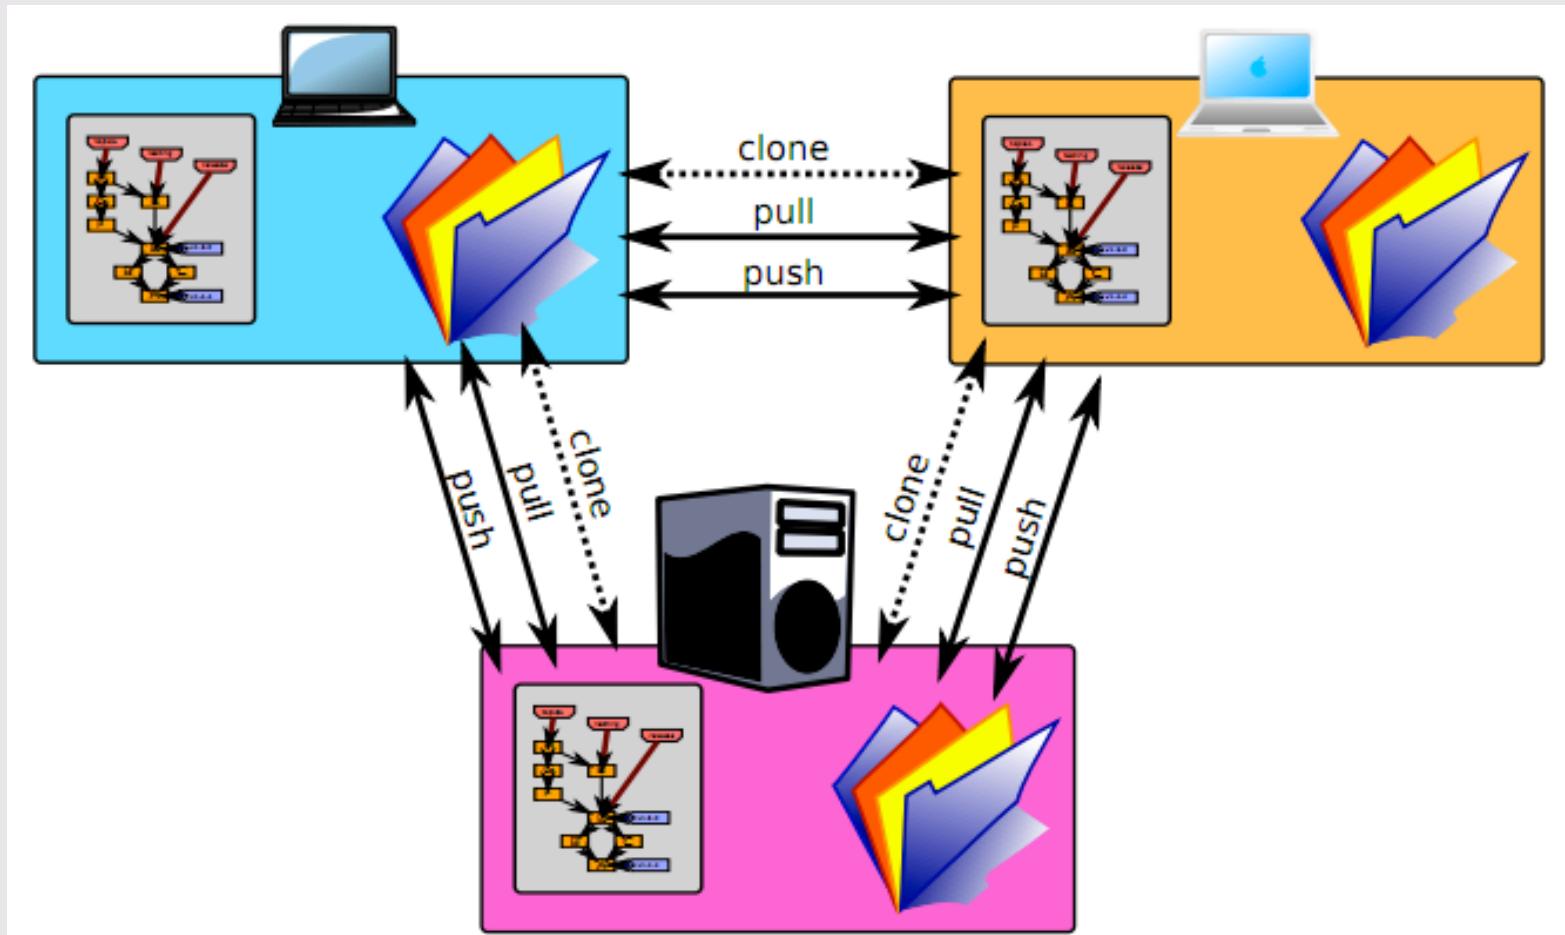
Problems

- Operations require a **server**
 - Single point of failure
 - Bottleneck



Git Primer – Decentralized SCM

- Anyone can be a server!



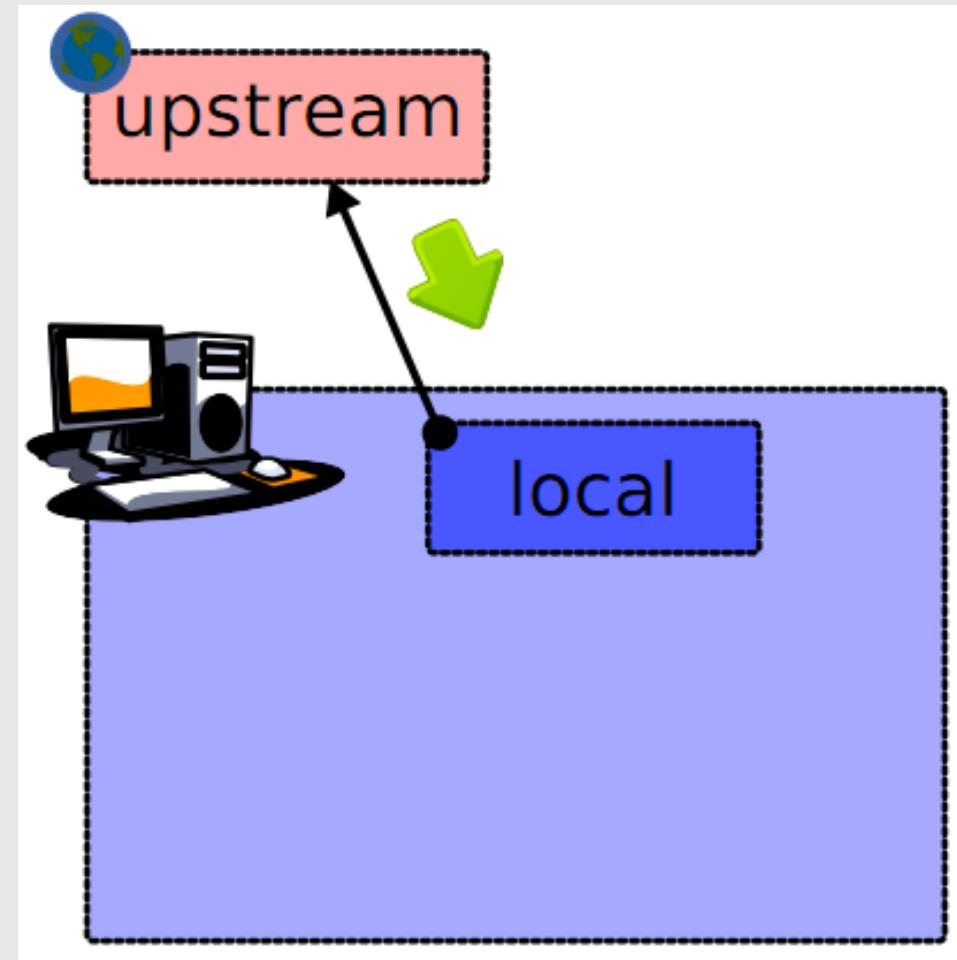
Git Primer - Decentralization

- Public Repository



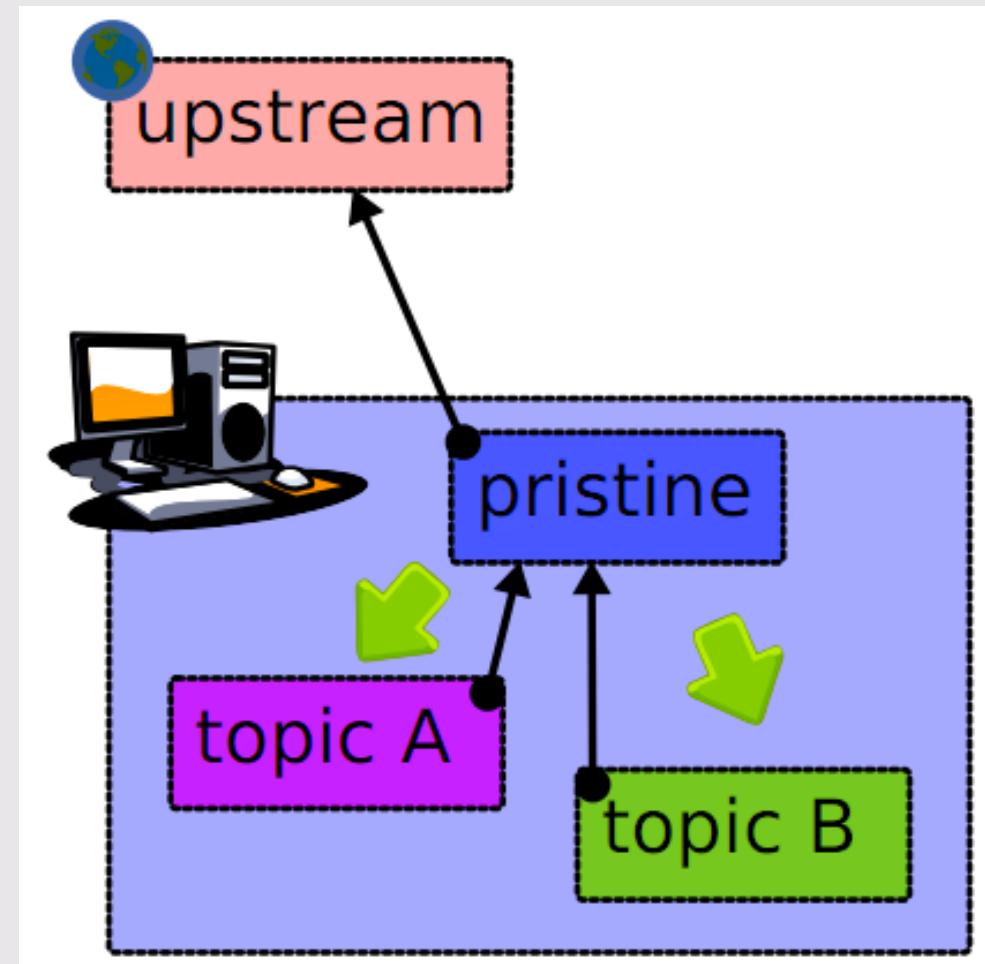
Git Primer - Decentralization

- Make a local clone



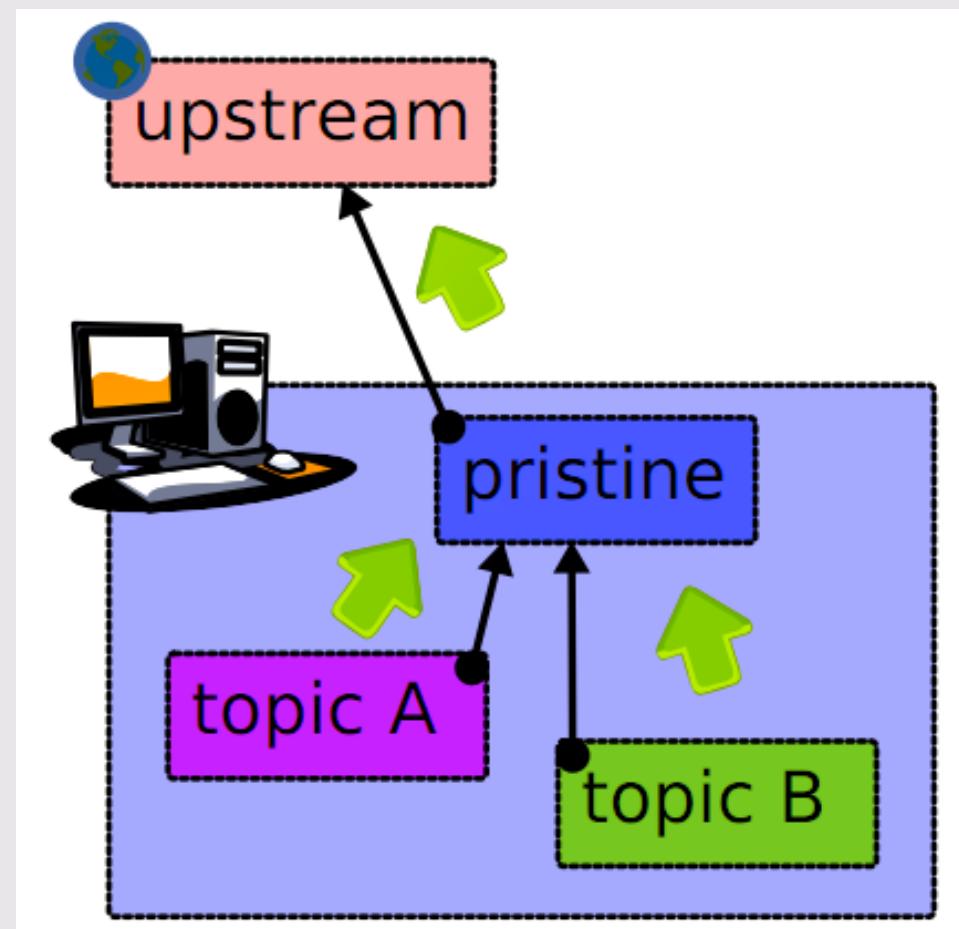
Git Primer - Decentralization

- Local cloning is lightweight



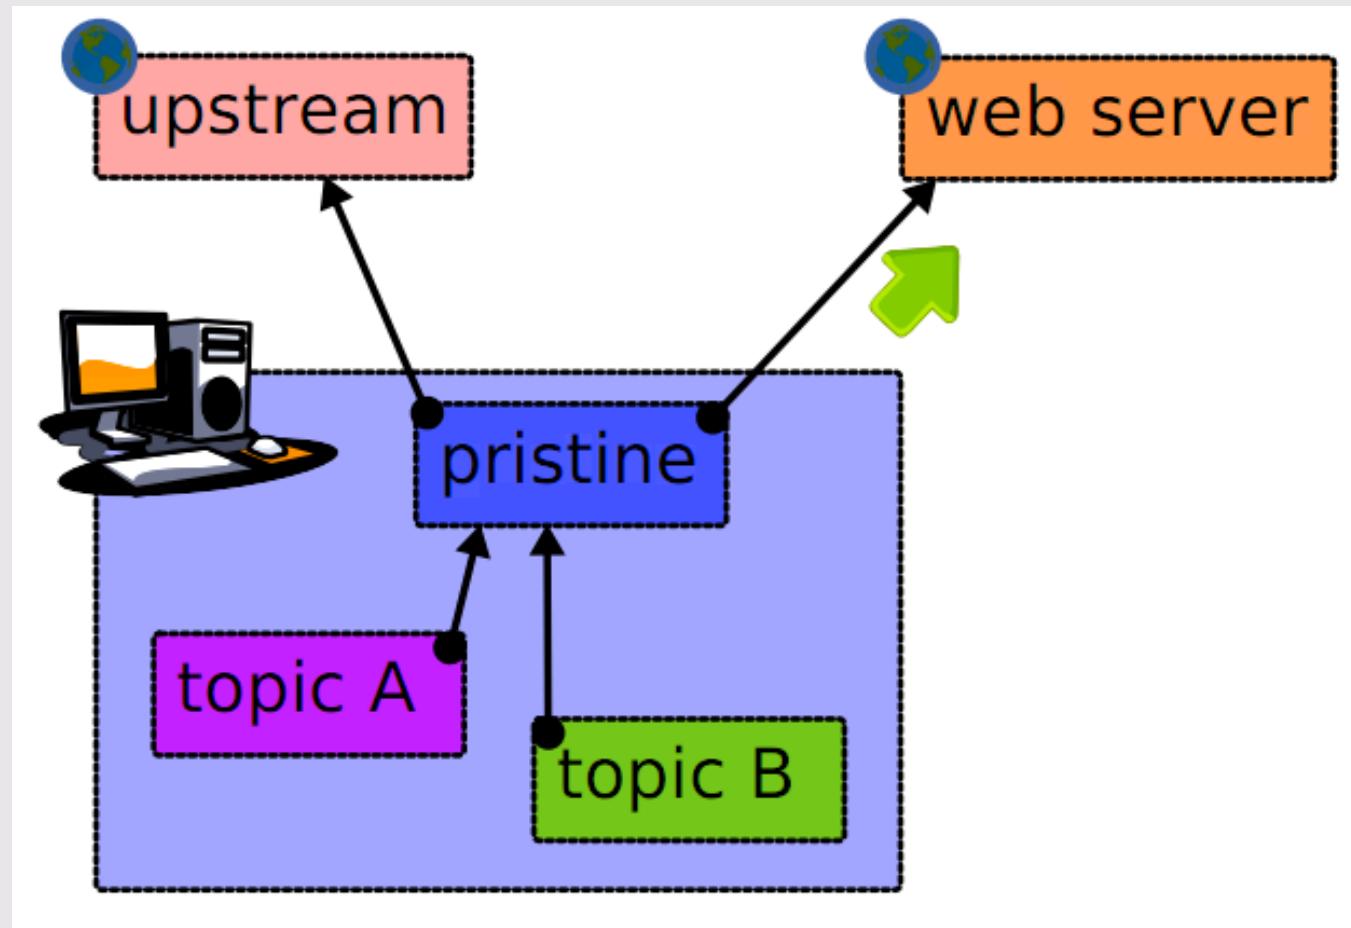
Git Primer - Decentralization

- Push changes between any repositories



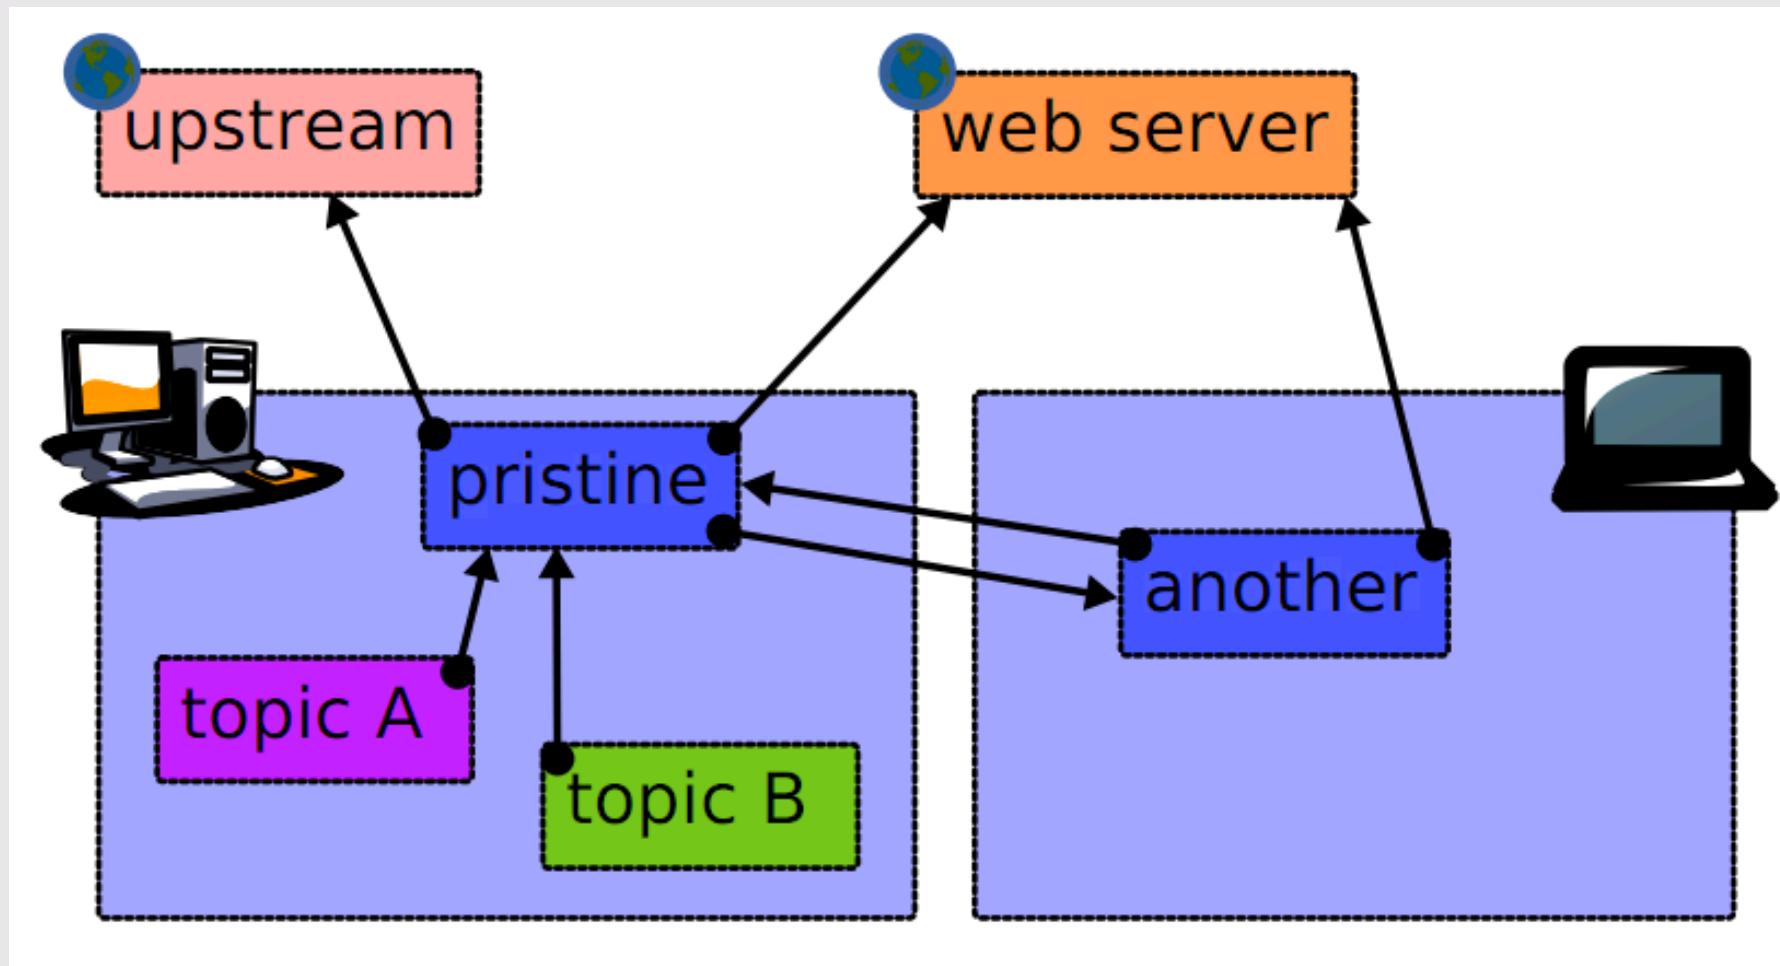
Git Primer - Decentralization

- Push changes to public server



Git Primer - Decentralization

- Share changes with trusted peers



But What About Perforce?

- Perforce is *the* SCM for Netflix and won't be changing anytime soon
- And your point is?
- Perforce is:
 - A server repository
 - A local repository
 - A schema for our code



Perforce

- Good server repository – deeply ingrained at Netflix
- Really bad local repository
 - Must be connected to server
 - Must checkout files
 - Cumbersome branching
 - Difficult to keep in sync with server
 - Easy to interfere with others' development
 - Bizarre client \leftrightarrow view mapping paradigm



Branching in Perforce – Step 1*

Ensure that you are able to submit files to the new codeline

Make sure you have a "mapping" for the new //depot/rell files in your client view. You do not necessarily need a mapping for the files you are branching from, but you do need a mapping for the files you are branching to, exactly as if you were adding new files to the depot. Your client spec should include a view mapping similar to the following:

```
//depot/rell/... //bruno_ws/rell/...
```

For example, a client spec that maps both codelines might look like:

```
Client: bruno_ws  
Root:   c:\workspace  
View:
```

```
//depot/main/... //bruno_ws/main/...  
//depot/rell/... //bruno_ws/rell/...
```



*From <http://kb.perforce.com/?article=4>

Branching in Perforce – Step 1 (cont'd)

In either case, //depot/re11 maps to c:\workspace\re11.

For more information on creating and updating client workspaces, refer to the P4 User's Guide.

NOTE: You can only branch or add files into depot directories where you have write permissions or greater. If you need to branch files to directories where you do not currently have at least write access, update the protections table. For more information on adjusting protections, consult with your Perforce administrator, or see the Perforce System Administrator's Guide chapter on Protections.



Branching in Perforce – Step 2

Create a branch specification

Create a branch spec to define the relationship between the mainline and the release branch. Branch specs are optional, but useful because they act as "shortcuts" for frequently used integration mappings. The branch spec for this example is called "rell" and is created with the following command:

```
p4 branch rell
```

The p4 branch command brings up a form to edit, similar to the p4 client command. In your text editor, you make the "View" section show that the //depot/main files is mapped to //depot/rell files whenever this branch spec is used:

```
Branch: rell
View:
  //depot/main/...  //depot/rell/...
```

Once this branch spec has been created, it can be used as a pre-formed mapping for any p4 integrate command between these two codelines.



Branching in Perforce – Step 3

Integrate using the branch specification

Use the branch spec to tell the **p4 integrate** command what to do. For example:

```
p4 integrate -b rel1
```

p4 integrate copies all the files from //depot/main into your client workspace and opens them for branch into //depot/rel1. When you submit the opened files, the release branch files are created in the depot, and Perforce records the fact that the files in //depot/rel1 came from files at the now-current revisions in //depot/main. The output from **p4 integrate** looks like:

```
//depot/rel1/a.c#1 - branch/sync from //depot/main/a.c#1  
//depot/rel1/b.c#1 - branch/sync from //depot/main/b.c#1
```

If you skipped step 2 and did not create a branch spec, the equivalent command would be:

```
p4 integrate //depot/main/... //depot/rel1/...
```

The output and final result is the same in either case.



Branching in Perforce – Step 4

Submit your branch

```
p4 submit
```

This brings up a submit form in your text editor, exactly as when submitting any other change to the depot. Fill out the description, save the file, and exit the editor.

The new //depot/rell codeline has now been submitted to the depot, synced to your workspace, and is ready to use.



Branching in Git

```
> git branch <branch name>
```

Um, that's it!



Switching between branches

```
> git checkout <branch name>
```

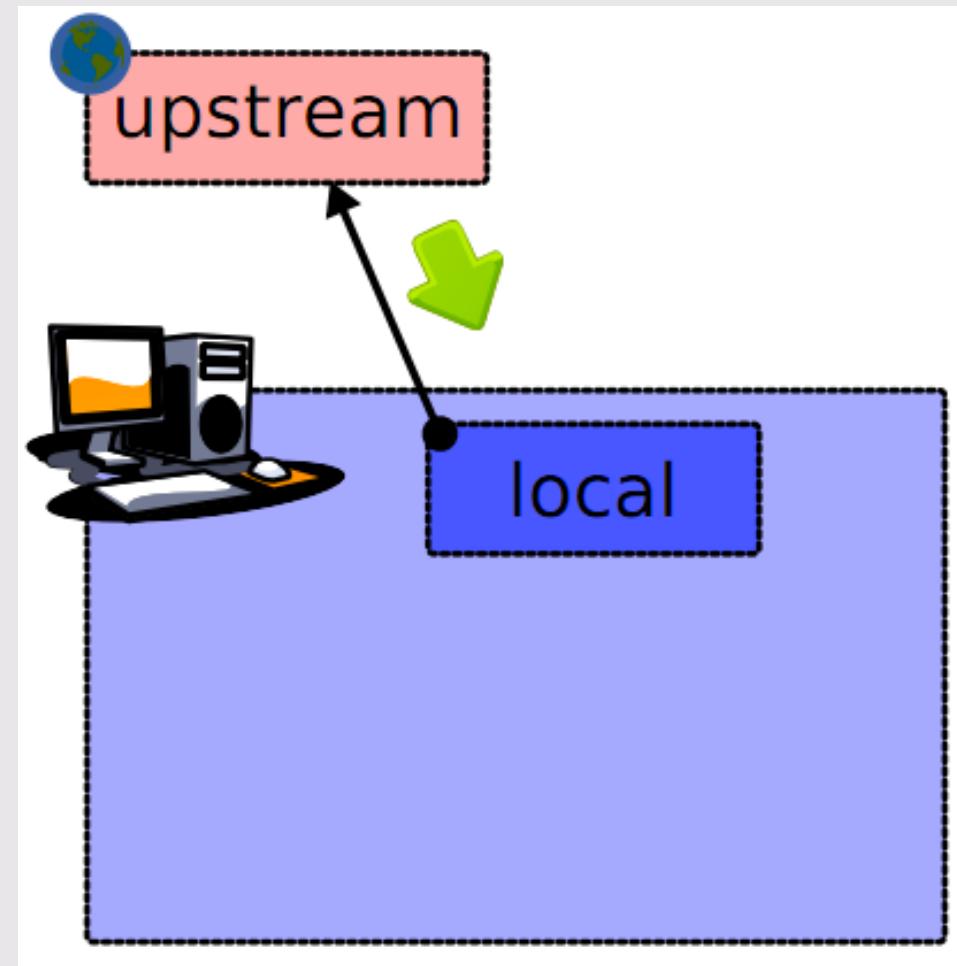
Super fast



Git and Perforce Together

- Remember this picture?

Perforce is here →



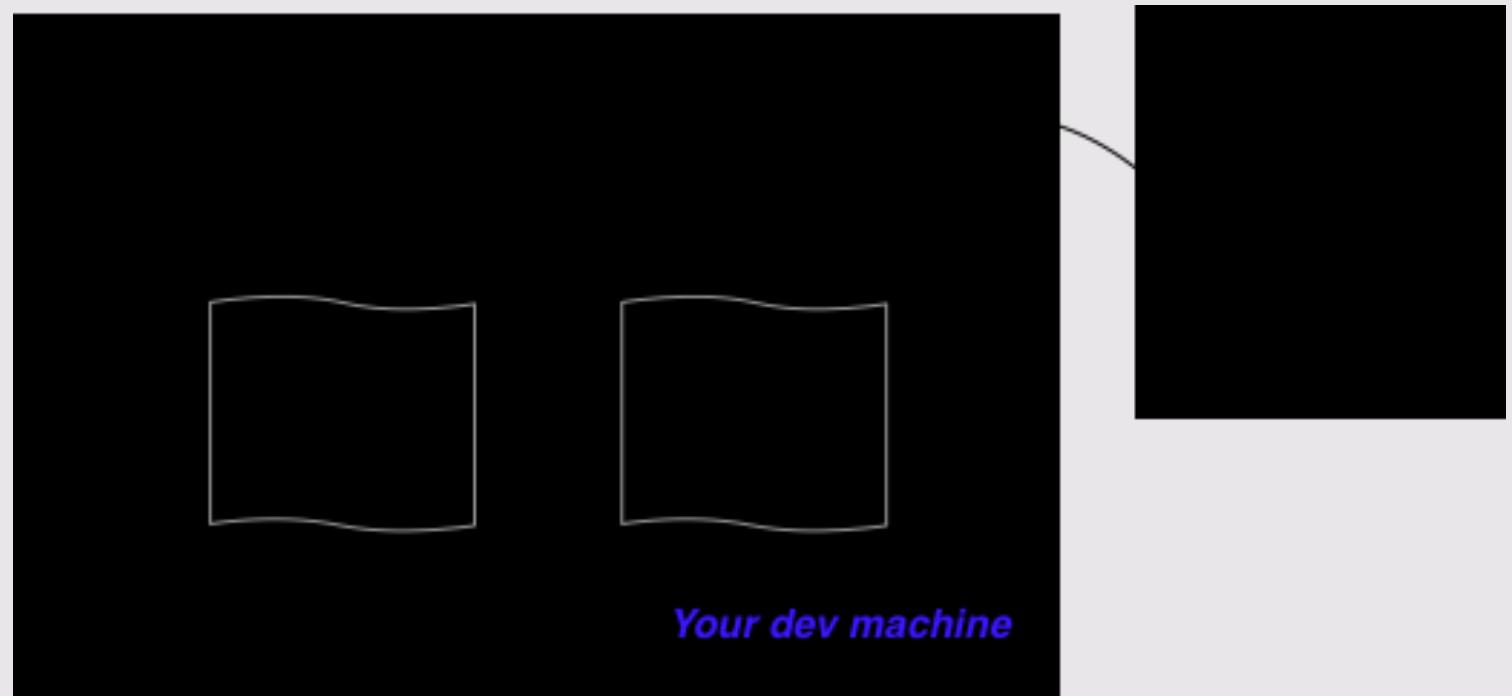
git-p4

- A Python script
- My fork with some changes:
<https://github.com/Randgalt/git-p4>
- A hack
- Replays git commits to Perforce
- I've been using it for months – it works great 99% of the time
 - The 1% when it doesn't can be frustrating



git-p4

- Keep Perforce workspace exactly the same
- A new git dir/local repo mirrors workspace
- git-p4 keeps the two in sync



Creating the Git Repo

```
> git-p4 clone //depot/commonlibraries/whatever
```

- You now have a git repo
- This is your main work folder (forget that the Perforce workspace exists)
- Commit, branch, etc. at will!
- When you're ready, sync back to Perforce:

```
> git-p4 submit --auto
```

- To sync changes *from* Perforce to git:

```
> git-p4 rebase
```



git-p4 Gets Confused Sometimes

- git-p4 stores context in your commit comments:
 - [git-p4: depot-paths = "//depot/commonlibraries/astyanax/": change = 928264]
- You really don't want to mess with the Perforce side of things
- It doesn't seem possible to sync git branches to Perforce



Git Tips/Tricks/Quirks

- Git is very easy to use – if you think Unix/Linux is easy to use ☺
- New concepts:
 - No checkout
 - Staging area
 - Local repo
 - No trunk
 - Make frequent commits/checkins



CLI Oddities

- checkout doesn't mean what you think it does
- Many commands require odd options:
 - reset --hard
 - rm --cached
 - Sometimes need to put files after --
 - git diff HEAD^ -- file/hash
- Just treat them as incantations – search the net to find out what to do
- Every CLI command accepts --help



No Checkout

- You don't need to tell git about files you're editing/adding/moving/deleting. Git will figure it out.

DEMO!



Staging Area

- Git calls this the “Index”
- A middle repo for staging potential commits
- Very powerful

DEMO!



Local Repo

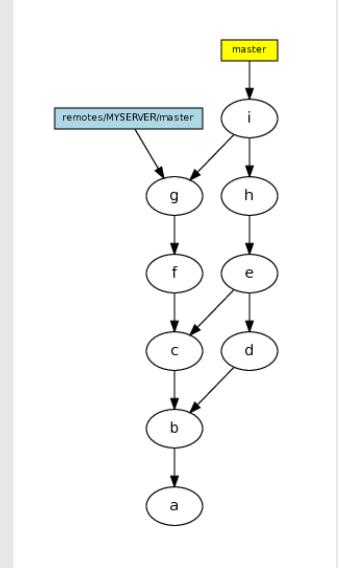
- Say it three times: *there is no server*
- Your repo is local, it's yours, no one else's
- Git is a distributed set of local repos that can be synced

DEMO!



No Trunk (*sorta*)

- A git repo is a DAG
- The head can be any node
- You can move from node to node easily
- In practice, there is always a master
- Many people flatten the DAG periodically



Ignoring Files

- You will need to duplicate any file patterns you have in `ignore` to `gitignore`

	<i>Perforce</i>	<u>Git</u>
<i>Single File</i>	<u>.ivy2.cache</u>	<u>.ivy2.cache</u>
<i>Directory</i>	<u>coverage</u>	<u>coverage/</u>
<i>File Extension</i>	<u>.foo</u>	<u>*.foo</u>



- Several ways to ignore in Git
 - `~/.gitignore`

git-p4

- The three important git-p4 commands:

git-p4 clone

Create a new repo



git-p4 submit --auto

Submit/Sync to Perforce



git-p4 rebase

Sync from Perforce



DEMO!



You should make bash shortcuts for these

How I Do Things

As an Example/Suggestion

- One dir per sub-project
- Git dir is completely separate from Perforce dir

```
~/Netflix/dev/projects$ ls -l | grep ^d
drwxr-xr-x 17 jzimmerman jzimmerman 578 Jun 25 11:52 apache-config
drwxr-xr-x  9 jzimmerman jzimmerman 306 Jul 19 15:29 astyanax
drwxr-xr-x 12 jzimmerman jzimmerman 408 Apr 19 16:42 async-http-client
drwxr-xr-x  3 jzimmerman jzimmerman 102 Apr 12 14:33 build
drwxr-xr-x  8 jzimmerman jzimmerman 272 Jul 12 17:12 cassandra-performance-test
drwxr-xr-x  9 jzimmerman jzimmerman 306 Jul 19 14:28 cassandra-test-harness
drwxr-xr-x  7 jzimmerman jzimmerman 238 Apr 25 15:45 cassandra_tools
drwxr-xr-x  3 jzimmerman jzimmerman 102 Jun 13 10:59 commons
drwxr-xr-x  7 jzimmerman jzimmerman 238 Jun 14 11:20 concurrent
drwxr-xr-x  9 jzimmerman jzimmerman 306 Jul  7 12:52 discovery
drwxr-xr-x  6 jzimmerman jzimmerman 204 Jul 19 13:08 epic-common
drwxr-xr-x  9 jzimmerman jzimmerman 306 Jul 19 12:55 generic-test-harness
drwxr-xr-x  7 jzimmerman jzimmerman 238 Jun  2 09:21 generic-test-harness-spi
drwxr-xr-x  8 jzimmerman jzimmerman 272 Jun 21 09:45 git-p4
drwxr-xr-x  5 jzimmerman jzimmerman 170 May 19 15:20 kway
drwxr-xr-x  5 jzimmerman jzimmerman 170 Jun  3 15:48 main
drwxr-xr-x  9 jzimmerman jzimmerman 306 Jun 13 11:52 netflixivy
drwxr-xr-x 11 jzimmerman jzimmerman 374 Jun  8 10:10 nfcassandra
drwxr-xr-x  9 jzimmerman jzimmerman 306 Jun 24 13:53 nflibrary
drwxr-xr-x  2 jzimmerman jzimmerman 68 Apr 12 15:56 paging
drwxr-xr-x  7 jzimmerman jzimmerman 238 Jul  8 16:42 platform
drwxr-xr-x  7 jzimmerman jzimmerman 238 May 19 13:44 platform-core
drwxr-xr-x 15 jzimmerman jzimmerman 510 Apr 13 12:47 post2crucible
drwxr-xr-x  7 jzimmerman jzimmerman 238 Jul 19 16:57 queue
drwxr-xr-x 10 jzimmerman jzimmerman 340 Jun 28 15:35 review-git
drwxr-xr-x  4 jzimmerman jzimmerman 136 May 31 14:44 statistics
drwxr-xr-x  9 jzimmerman jzimmerman 306 Jul 14 11:42 zookeeper-client
```



Branching

DEMO!



We are starting a Netflix presence on github!

Github

github
SOCIAL CODING

Randgalt | Dashboard | Inbox 0 | Account Settings | Log Out

Explore GitHub Gist Blog Help Search...

Netflix / curator

Admin Unwatch Fork Your Fork Pull Request 2 4 2

Source Commits Network Pull Requests (0) Fork Queue Issues (0) Wiki (4) Graphs Branch: master

Switch Branches (1) ▾ Switch Tags (0) Branch List Search source code...

ZooKeeper client wrapper and rich ZooKeeper framework — [Read more](#)
[click here to add a homepage](#)

Downloads

Fork anyone's project
(even forks of forks of
...)

Make changes and then request to have
them merged back to the parent project



SO much better than SVN
hosting!

Patches

DEMO!



Stash

DEMO!



Code Reviews

- I've written a CLI app to add Code Reviews to Crucible from Git
 - review-git.jar
 - <https://confluence.corp.netflix.com/display/ENGTOOLS/Git+with+Perforce>



GUIs/IDEs

- Git integrates with most IDEs
- I use the CLI `~/Netflix/dev/projects/astyanax[async*]$/`
 - There are some great bash prompt scripts
- For Mac: gitx
 - <http://gitx.frim.nl/>
- For Windows: TortoiseGIT
 - <http://code.google.com/p/tortoisegit/>
- Linux: giggle
 - <http://live.gnome.org/giggle>

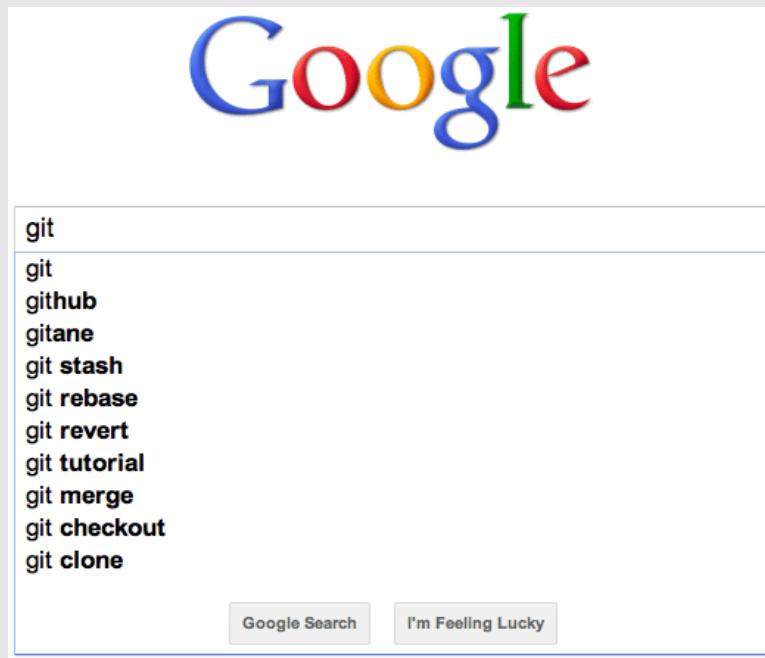
Branch name

Means uncommitted changes



So, So, Much More...

- I've only scratched the surface
- <http://git-scm.com/>
- Everyday Git
 - <http://www.kernel.org/pub/software/scm/git/docs/everyday.html>
-



Thank you for coming!

Questions

